

Function List and Descriptions

`prepare_and_split_corpus(ds)`

- Concatenates 'train', 'validation', and 'test' dataset text into a single corpus, then splits it into training, validation, and test sets by 80%, 10% and 10%.

`tokenize(text)`

- Performs a basic whitespace tokenization on the input text. Used to prepare text for custom BPE tokenizer training.

`compute_coverage(tokenizer, dataset, sample_size=1000)`

- Calculates the percentage of tokens in the dataset that are covered by the tokenizer (i.e., not mapped to <UNK>). Useful for evaluating tokenizer quality.

`encode_text_str(tokenizer, text)`

- Encodes raw text into a list of token IDs using the trained BPE tokenizer.

`build_input_target_pairs(token_ids, block_size=32)`

- Generates (input, target) pairs using a sliding window approach for training next-token prediction models.

`LanguageModelingDataset`

- A custom PyTorch Dataset class that wraps tokenized input-target sequences for use with DataLoader (shuffle the data in each epoch).

`EmbeddingLayer`

- Combines token and positional embeddings into dense vectors with dropout. Forms the input layer of the GPT model.

`SelfAttention`

- Implements multi-head self-attention using dot-product attention mechanism to capture inter-token dependencies.

`FeedForward`

- A two-layer MLP with ReLU activation and dropout, applied independently to each token position.

`TransformerBlock`

- A single transformer block composed of LayerNorm, multi-head attention, and feed-forward layers with residual connections.

`GPT`

- The main transformer-based language model class. Stacks multiple TransformerBlocks with final LayerNorm and linear output head.

`compute_loss(model, batch)`

- Computes cross-entropy loss between model outputs and target token

sequences.

`train_one_epoch(model, dataloader, optimizer, loss_fn, device)`
- Trains the model for one epoch by running forward and backward passes, and updating weights.

`evaluate_loss(model, val_loader, loss_fn, device)`
- Evaluates average loss on the validation set.

`evaluate_perplexity(model, dataloader)`
- Calculates perplexity from average validation loss. Perplexity is an evaluation metric for language models.

`compute_grad_norm(model)`
- Computes the total gradient norm across model parameters to monitor training stability.

`hyperparameter_search(...)`
- Performs a grid search over model hyperparameters such as learning rate, number of layers/heads, embedding dimension and returns the best configuration.
- Use validation data to fine tune
- Each config runs 10 epochs

`retrain_with_best_parameters(...)`
- Perform a retrain on the tuned best parameters
- Use validation data to fine tune
- Use on test data to compute perplexity
- Each config runs 20 epochs